

DL_Practical_4_Compact

November 12, 2025

Assignment No: 04

Aim: To implement Autoencoder for anomaly detection.

Problem Statement: Use Autoencoder to implement anomaly detection. Build the model by using:
a. Import required libraries b. Upload / access the dataset c. Encoder converts it into latent representation d. Decoder networks convert it back to the original input e. Compile the models with Optimizer, Loss, and Evaluation Metrics. Objectives: a) Apply Autoencoder deep learning architecture to determine anomalies in input dataset b) Evaluate Model.

```
[1]: # a) Import required libraries

import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers, Model, optimizers, losses, metrics

print("Imports complete!")

# b) Upload / access the dataset

with np.load(r"C:\Users\kusha\Desktop\mnist_dataset.npz") as data:
    x_train_all = data["X_train"]
    y_train_all = data["y_train"]
    x_test_all = data["X_test"]
    y_test_all = data["y_test"]

    print(
        f"x_train_all:\n"
        f"  Data type: {x_train_all.dtype}\n"
        f"  Shape     : {x_train_all.shape}\n"
        f"  Pixel value range: {x_train_all.min()} to {x_train_all.max()}\n"
    )
    print(
        f"y_train_all:\n"
        f"  Data type: {y_train_all.dtype}\n"
        f"  Shape     : {y_train_all.shape}\n"
        f"  First 10 labels: {y_train_all[:10]}\n"
    )
```

```

print(
    f"x_test_all:\n"
    f" Data type: {x_test_all.dtype}\n"
    f" Shape : {x_test_all.shape}\n"
    f" Pixel value range: {x_test_all.min()} to {x_test_all.max()}\n"
)
print(
    f"y_test_all:\n"
    f" Data type: {y_test_all.dtype}\n"
    f" Shape : {y_test_all.shape}\n"
    f" First 10 labels: {y_test_all[:10]}\n"
)
print(f"First image sample of x_train_all (pixel values):\n{x_train_all[0]}\n")

# Normalize and reshape for CNN
x_train_all = (x_train_all.astype("float32") / 255.0)[..., None]
x_test_all = (x_test_all.astype("float32")) / 255.0[..., None]

# Use digit '1' as normal class for training (anomaly detection)
normal_class = 1
x_train = x_train_all[y_train_all == normal_class]
x_test = x_test_all
y_test = y_test_all

print(
    f"\nAfter normalization and selection for anomaly detection:\n"
    f"Training on digit '{normal_class}' only.\n"
    f"x_train: {x_train.shape}, x_test: {x_test.shape}\n"
)
# c) Encoder network: converts input to latent representation

input_shape = (28, 28, 1)
latent_dim = 16

encoder_inputs = layers.Input(shape=input_shape)
x = layers.Conv2D(32, 3, strides=2, padding='same', ↴
    activation='relu')(encoder_inputs) # (14, 14, 32)
x = layers.Conv2D(64, 3, strides=2, padding='same', activation='relu')(x) ↴
    # (7, 7, 64)
x = layers.Flatten()(x)
latent = layers.Dense(latent_dim, name="latent")(x)
encoder = Model(encoder_inputs, latent, name="encoder")

print("Encoder architecture:")
encoder.summary()

```

```

# d) Decoder network: reconstructs input from latent

latent_inputs = layers.Input(shape=(latent_dim,))
x = layers.Dense(7*7*64, activation='relu')(latent_inputs)
x = layers.Reshape((7,7,64))(x)
x = layers.Conv2DTranspose(64, 3, strides=2, padding='same',
    activation='relu')(x) # (14,14,64)
x = layers.Conv2DTranspose(32, 3, strides=2, padding='same',
    activation='relu')(x) # (28,28,32)
decoded = layers.Conv2D(1, 3, padding='same', activation='sigmoid')(x)
    # (28,28,1)
decoder = Model(latent_inputs, decoded, name="decoder")

print("Decoder architecture:")
decoder.summary()

# e) Compile the complete autoencoder

autoencoder_inputs = layers.Input(shape=input_shape)
encoded = encoder(autoencoder_inputs)
reconstructed = decoder(encoded)
autoencoder = Model(autoencoder_inputs, reconstructed, name="autoencoder")

autoencoder.compile(
    optimizer=optimizers.Adam(learning_rate=0.001),
    loss=losses.BinaryCrossentropy(),
    metrics=[metrics.MeanSquaredError()]
)

print("Autoencoder architecture:")
autoencoder.summary()

# Train the model (train only on normal class)

history = autoencoder.fit(
    x_train, x_train,
    epochs=20,
    batch_size=128,
    validation_split=0.1,
    verbose=2
)
print("Training complete!\n")

# Evaluate/analyze: Compute reconstruction error threshold for anomaly detection

# Reconstruct training data for threshold
recon_train = autoencoder.predict(x_train)

```

```

train_mse = np.mean((recon_train - x_train) ** 2, axis=(1,2,3))

# Use simple threshold: mean + 3*std of normal training errors
threshold = train_mse.mean() + 3 * train_mse.std()
print(f"Anomaly threshold (mean + 3*std): {threshold:.6f}")

# Reconstruct test data
recon_test = autoencoder.predict(x_test)
test_mse = np.mean((recon_test - x_test) ** 2, axis=(1,2,3))
anomaly_flags = test_mse > threshold

print("Total test samples:", len(x_test))
print("Detected anomalies:", np.sum(anomaly_flags))

# Visualization: Show original and reconstruction, highlight anomalies vs. ↴normal

is_normal_test = (y_test == normal_class)
n = 6
plt.figure(figsize=(12, 5))
example_idx = np.concatenate([
    np.where(is_normal_test)[0][:n//2],
    np.where(~is_normal_test)[0][:n - n//2]
])

for i, idx in enumerate(example_idx):
    # original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[idx].reshape(28,28), cmap='gray', vmin=0, vmax=1)
    plt.title(f"True:{y_test[idx]} Err:{test_mse[idx]:.4f}")
    plt.axis('off')
    # reconstruction
    ax = plt.subplot(2, n, n + i + 1)
    plt.imshow(recon_test[idx].reshape(28,28), cmap='gray', vmin=0, vmax=1)
    flag = "ANOMALY" if anomaly_flags[idx] else "NORMAL"
    plt.title(flag)
    plt.axis('off')

plt.suptitle("Original (top row) vs Reconstruction (bottom row)")
plt.tight_layout(rect=[0, 0, 1, 0.95])
plt.show()

```

Imports complete!

x_train_all:

 Data type: uint8
 Shape : (60000, 28, 28)
 Pixel value range: 0 to 255

After normalization and selection for anomaly detection:

Training on digit '1' only.

x_train: (6742, 28, 28, 1), x_test: (10000, 28, 28, 1)

Encoder architecture:

Model: "encoder"

Layer (type)	Output Shape
input_layer (InputLayer) ↳ Param # 0	(None, 28, 28, 1)
conv2d (Conv2D) ↳ 320	(None, 14, 14, 32)
conv2d_1 (Conv2D) ↳ 18,496	(None, 7, 7, 64)

```
flatten (Flatten)           (None, 3136) ▾  
↳ 0
```

```
latent (Dense)             (None, 16) ▾  
↳ 50,192
```

Total params: 69,008 (269.56 KB)

Trainable params: 69,008 (269.56 KB)

Non-trainable params: 0 (0.00 B)

Decoder architecture:

Model: "decoder"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

input_layer_1 (InputLayer)	(None, 16)	0
----------------------------	------------	---

dense (Dense)	(None, 3136)	53,312
---------------	--------------	--------

reshape (Reshape)	(None, 7, 7, 64)	0
-------------------	------------------	---

conv2d_transpose (Conv2DTranspose)	(None, 14, 14, 64)	36,928
------------------------------------	--------------------	--------

conv2d_transpose_1 (Conv2DTranspose)	(None, 28, 28, 32)	18,464
--------------------------------------	--------------------	--------

conv2d_2 (Conv2D)	(None, 28, 28, 1)	289
-------------------	-------------------	-----

Total params: 108,993 (425.75 KB)

Trainable params: 108,993 (425.75 KB)

```
Non-trainable params: 0 (0.00 B)
```

```
Autoencoder architecture:
```

```
Model: "autoencoder"
```

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 28, 28, 1)	0
encoder (Functional)	(None, 16)	69,008
decoder (Functional)	(None, 28, 28, 1)	108,993

```
Total params: 178,001 (695.32 KB)
```

```
Trainable params: 178,001 (695.32 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
Epoch 1/20
```

```
48/48 - 16s - 337ms/step - loss: 0.3299 - mean_squared_error: 0.0957 - val_loss: 0.1357 - val_mean_squared_error: 0.0325
```

```
Epoch 2/20
```

```
48/48 - 5s - 108ms/step - loss: 0.1322 - mean_squared_error: 0.0314 - val_loss: 0.1168 - val_mean_squared_error: 0.0264
```

```
Epoch 3/20
```

```
48/48 - 5s - 109ms/step - loss: 0.0973 - mean_squared_error: 0.0197 - val_loss: 0.0757 - val_mean_squared_error: 0.0127
```

```
Epoch 4/20
```

```
48/48 - 5s - 104ms/step - loss: 0.0736 - mean_squared_error: 0.0120 - val_loss: 0.0631 - val_mean_squared_error: 0.0088
```

```
Epoch 5/20
```

```
48/48 - 5s - 108ms/step - loss: 0.0599 - mean_squared_error: 0.0077 - val_loss: 0.0558 - val_mean_squared_error: 0.0065
```

```
Epoch 6/20
```

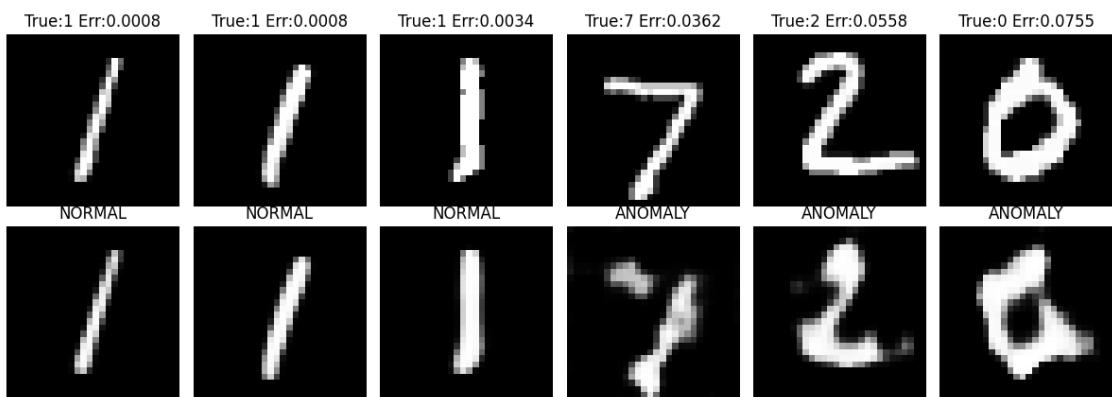
```
48/48 - 4s - 91ms/step - loss: 0.0539 - mean_squared_error: 0.0058 - val_loss: 0.0504 - val_mean_squared_error: 0.0048
```

```
Epoch 7/20
```

```
48/48 - 6s - 127ms/step - loss: 0.0495 - mean_squared_error: 0.0045 - val_loss:  
0.0472 - val_mean_squared_error: 0.0038  
Epoch 8/20  
48/48 - 5s - 109ms/step - loss: 0.0470 - mean_squared_error: 0.0038 - val_loss:  
0.0457 - val_mean_squared_error: 0.0034  
Epoch 9/20  
48/48 - 6s - 120ms/step - loss: 0.0455 - mean_squared_error: 0.0033 - val_loss:  
0.0444 - val_mean_squared_error: 0.0030  
Epoch 10/20  
48/48 - 5s - 99ms/step - loss: 0.0443 - mean_squared_error: 0.0030 - val_loss:  
0.0435 - val_mean_squared_error: 0.0028  
Epoch 11/20  
48/48 - 5s - 97ms/step - loss: 0.0434 - mean_squared_error: 0.0027 - val_loss:  
0.0428 - val_mean_squared_error: 0.0026  
Epoch 12/20  
48/48 - 4s - 91ms/step - loss: 0.0428 - mean_squared_error: 0.0025 - val_loss:  
0.0424 - val_mean_squared_error: 0.0025  
Epoch 13/20  
48/48 - 4s - 93ms/step - loss: 0.0423 - mean_squared_error: 0.0024 - val_loss:  
0.0419 - val_mean_squared_error: 0.0023  
Epoch 14/20  
48/48 - 5s - 97ms/step - loss: 0.0419 - mean_squared_error: 0.0023 - val_loss:  
0.0416 - val_mean_squared_error: 0.0022  
Epoch 15/20  
48/48 - 5s - 96ms/step - loss: 0.0415 - mean_squared_error: 0.0022 - val_loss:  
0.0414 - val_mean_squared_error: 0.0022  
Epoch 16/20  
48/48 - 5s - 112ms/step - loss: 0.0413 - mean_squared_error: 0.0021 - val_loss:  
0.0411 - val_mean_squared_error: 0.0021  
Epoch 17/20  
48/48 - 5s - 94ms/step - loss: 0.0410 - mean_squared_error: 0.0020 - val_loss:  
0.0409 - val_mean_squared_error: 0.0020  
Epoch 18/20  
48/48 - 5s - 102ms/step - loss: 0.0408 - mean_squared_error: 0.0020 - val_loss:  
0.0408 - val_mean_squared_error: 0.0020  
Epoch 19/20  
48/48 - 4s - 93ms/step - loss: 0.0406 - mean_squared_error: 0.0019 - val_loss:  
0.0405 - val_mean_squared_error: 0.0019  
Epoch 20/20  
48/48 - 4s - 86ms/step - loss: 0.0403 - mean_squared_error: 0.0018 - val_loss:  
0.0405 - val_mean_squared_error: 0.0019  
Training complete!
```

```
211/211          5s 20ms/step  
Anomaly threshold (mean + 3*std): 0.006872  
313/313          6s 18ms/step  
Total test samples: 10000  
Detected anomalies: 8891
```

Original (top row) vs Reconstruction (bottom row)



[]: