

Assignment No: 06 Aim: Implementation of object detection using transfer learning of CNN architectures Problem Statement:
Implementation of object detection using transfer learning of CNN architecture. a. Load in a pre-trained CNN model trained on a large dataset b. Freeze parameters (weights) in model's lower convolutional layers c. Add custom classifier with several layers of trainable parameters to model d. Train classifier layers on training data available for task e. Fine-tune hyperparameters and unfreeze more layers as needed

```
In [1]: # a) Load in a pre-trained CNN model trained on a large dataset
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.applications import VGG16
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

with np.load(r"C:\Users\kusha\Desktop\mnist_dataset.npz") as data:
    x_train = data["X_train"]
    y_train = data["y_train"]
    x_test = data["X_test"]
    y_test = data["y_test"]

print("✅ Loaded MNIST locally!")
print("Training data shape:", x_train.shape)
print("Test data shape:", x_test.shape)

# b) Freeze parameters (weights) in model's Lower convolutional Layers
x_train = np.stack((x_train,) * 3, axis=-1)
x_test = np.stack((x_test,) * 3, axis=-1)

x_train = tf.image.resize(x_train, [32, 32])
x_test = tf.image.resize(x_test, [32, 32])

x_train = x_train / 255.0
x_test = x_test / 255.0

y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
for layer in base_model.layers:
    layer.trainable = False
print("Frozen base model layers.")

# c) Add custom classifier with trainable Layers
model = models.Sequential([
    base_model,
    layers.Flatten(),
    layers.Dense(256, activation='relu'),
    layers.Dropout(0.4),
    layers.Dense(10, activation='softmax') # 10 classes in MNIST
])

# d) Train classifier layers on training data
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

print("Starting training...")
model.summary()
```

```

history = model.fit(
    x_train, y_train,
    epochs=5,
    batch_size=64,
    validation_split=0.1,
    verbose=1
)
print("Initial training complete.")

# e) Fine-tune hyperparameters and unfreeze more layers
for layer in base_model.layers[-4:]:
    layer.trainable = True
print("Unfroze last 4 layers of base model for fine-tuning.")

model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

fine_tune_history = model.fit(
    x_train, y_train,
    epochs=3,
    batch_size=64,
    validation_split=0.1,
    verbose=1
)

print("Fine-tuning complete.")

# f) Evaluate and plot results
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
print(f"Test Accuracy after fine-tuning: {test_acc:.4f}")
print(f"Test Loss after fine-tuning: {test_loss:.4f}")

plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy (Base)')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy (Base)')
plt.plot(fine_tune_history.history['accuracy'], label='Train Accuracy (Fine-Tuned)')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss (Base)')
plt.plot(history.history['val_loss'], label='Validation Loss (Base)')
plt.plot(fine_tune_history.history['loss'], label='Train Loss (Fine-Tuned)')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()

```


✓ Loaded MNIST locally!
Training data shape: (60000, 28, 28)
Test data shape: (10000, 28, 28)
Froze base model layers.
Starting training...
Model: "sequential"

Layer (type)	Output Shape
vgg16 (Functional)	(None, 1, 1, 512)
flatten (Flatten)	(None, 512)
dense (Dense)	(None, 256)
dropout (Dropout)	(None, 256)
dense_1 (Dense)	(None, 10)



Total params: 14,848,586 (56.64 MB)
Trainable params: 133,898 (523.04 KB)
Non-trainable params: 14,714,688 (56.13 MB)

Epoch 1/5
844/844 94s 110ms/step - accuracy: 0.7497 - loss: 0.7942 - val_accuracy: 0.9437 - val_loss: 0.1802
Epoch 2/5
844/844 97s 115ms/step - accuracy: 0.9292 - loss: 0.2279 - val_accuracy: 0.9605 - val_loss: 0.1263
Epoch 3/5
844/844 102s 121ms/step - accuracy: 0.9479 - loss: 0.1688 - val_accuracy: 0.9640 - val_loss: 0.1067
Epoch 4/5
844/844 96s 114ms/step - accuracy: 0.9547 - loss: 0.1443 - val_accuracy: 0.9635 - val_loss: 0.1086
Epoch 5/5
844/844 97s 115ms/step - accuracy: 0.9568 - loss: 0.1323 - val_accuracy: 0.9682 - val_loss: 0.0901
Initial training complete.
Unfroze last 4 layers of base model for fine-tuning.
Epoch 1/3
844/844 399s 470ms/step - accuracy: 0.9654 - loss: 0.1064 - val_accuracy: 0.9842 - val_loss: 0.0487
Epoch 2/3
844/844 366s 434ms/step - accuracy: 0.9842 - loss: 0.0493 - val_accuracy: 0.9867 - val_loss: 0.0428
Epoch 3/3
844/844 337s 400ms/step - accuracy: 0.9888 - loss: 0.0356 - val_accuracy: 0.9910 - val_loss: 0.0283
Fine-tuning complete.
Test Accuracy after fine-tuning: 0.9872
Test Loss after fine-tuning: 0.0423




```
In [4]: # Predict on test set
predictions = model.predict(x_test)
pred_labels = np.argmax(predictions, axis=1)
true_labels = np.argmax(y_test, axis=1)

# Show first 10 test images with predicted and true labels
plt.figure(figsize=(15, 3))
for i in range(10):
    plt.subplot(1, 10, i+1)
    plt.imshow(x_test[i])
    plt.title(f"P: {pred_labels[i]}\nT: {true_labels[i]}")
    plt.axis('off')
plt.suptitle("Predicted vs True Labels on test images")
plt.show()
```

313/313 ━━━━━━━━ 18s 58ms/step

Predicted vs True Labels on test images



In []: